

Exercise 5: Solution

Non-linearities

Sigmoid - Forward

```
def forward(self, x):
    """
    :param x: Inputs, of any shape

    :return out: Output, of the same shape as x
    :return cache: Cache, for backward computation, of the same shape as x
    """
    shape = x.shape
    outputs, cache = np.zeros(shape), np.zeros(shape)
    #####
    # TODO:                                     #
    # Implement the forward pass of Sigmoid activation function #
    #####
    outputs = 1 / (1 + np.exp(-x))
    cache = outputs
    #####
    #                                     END OF YOUR CODE                                     #
    #####
    return outputs, cache
```

Remark:

The output of sigmoid function is stored in the cache for the computation in backward pass.

Sigmoid - Backward

```
def backward(self, dout, cache):
    """
    :return: dx: the gradient w.r.t. input X, of the same shape as X
    """
    dx = None
    #####
    # TODO:                                     #
    # Implement the backward pass of Sigmoid activation function #
    #####
    dx = dout * cache * (1 - cache)
    #####
    #                                     END OF YOUR CODE                                     #
    #####
    return dx
```

Remark:

The derivative of sigmoid function is
is $\text{sigmoid} * (1 - \text{sigmoid})$

Relu - Forward

```
def forward(self, x):
    """
    :param x: Inputs, of any shape

    :return out: Output, of the same shape as x
    :return cache: Cache, for backward computation, of the same shape as x
    """
    outputs = None
    cache = None

    #####
    # TODO:                                     #
    # Implement the forward pass of Relu activation function      #
    #####
    outputs = np.maximum(x, 0)
    cache = x

    #####
    #                                     END OF YOUR CODE      #
    #####
    return outputs, cache
```

Relu - Backward

```
def backward(self, dout, cache):
    """
    :return: dx: the gradient w.r.t. input X, of the same shape as X
    """

    dx = None

    #####
    # TODO:                                     #
    # Implement the backward pass of Relu activation function      #
    #####

    x = cache
    dx = dout
    dx[x < 0] = 0

    #####
    #                                     END OF YOUR CODE      #
    #####

    return dx
```

Affine Layers

Affine Layer- Forward

```
def affine_forward(x, w, b):
    """
    Computes the forward pass for an affine (fully-connected) layer.
    The input x has shape (N, d_1, ..., d_k) and contains a minibatch of N
    examples, where each example x[i] has shape (d_1, ..., d_k). We will
    reshape each input into a vector of dimension D = d_1 * ... * d_k, and
    then transform it to an output vector of dimension M.
    Inputs:
    :param x: A numpy array containing input data, of shape (N, d_1, ..., d_k)
    :param w: A numpy array of weights, of shape (D, M)
    :param b: A numpy array of biases, of shape (M,)
    :return out: output, of shape (N, M)
    :return cache: (x, w, b)
    """
    N, M = x.shape[0], b.shape[0]
    out = np.zeros((N,M))
    #####
    # TODO: Implement the affine forward pass. Store the result in out.  #
    # You will need to reshape the input into rows.                    #
    #####
    x_resaped = np.reshape(x, (x.shape[0], -1))
    out = x_resaped.dot(w) + b
    #####
    #                                END OF YOUR CODE                                #
    #####
    cache = (x, w, b)
    return out, cache
```

Remark: the input x , weights w and bias b are saved in cache, such that the backward pass can access them.

Affine Layer - Backward

```
def affine_backward(dout, cache):
    """
    Computes the backward pass for an affine layer.
    Inputs:
    :param dout: Upstream derivative, of shape (N, M)
    :param cache: Tuple of:
        - x: Input data, of shape (N, d_1, ... d_k)
        - w: Weights, of shape (D, M)
        - b: A numpy array of biases, of shape (M,)
    :return dx: Gradient with respect to x, of shape (N, d1, ..., d_k)
    :return dw: Gradient with respect to w, of shape (D, M)
    :return db: Gradient with respect to b, of shape (M,)
    """
    x, w, b = cache
    dx, dw, db = None, None, None
    #####
    # TODO: Implement the affine backward pass.
    # Hint: Don't forget to average the gradients dw and db
    #####
    n = x.shape[0]

    dw = (np.reshape(x, (x.shape[0], -1)).T).dot(dout) / n
    dw = np.reshape(dw, w.shape)

    db = np.mean(dout, axis=0, keepdims=False)

    dx = dout.dot(w.T)
    dx = np.reshape(dx, x.shape)
    #####
    #                               END OF YOUR CODE
    #####
    return dx, dw, db
```

Remark:

Make sure the dw and dx have the same shape as w and x .

Here, we take the average of the gradient, because otherwise we have the sum over all gradients $\sum_{i=1}^N \nabla L_i(\theta)$ of the entire minibatch.

Hint: This averaging operation can also be done in the backward pass of the loss function. If so, we don't average n in the layer.

Questions? Piazza 😊